

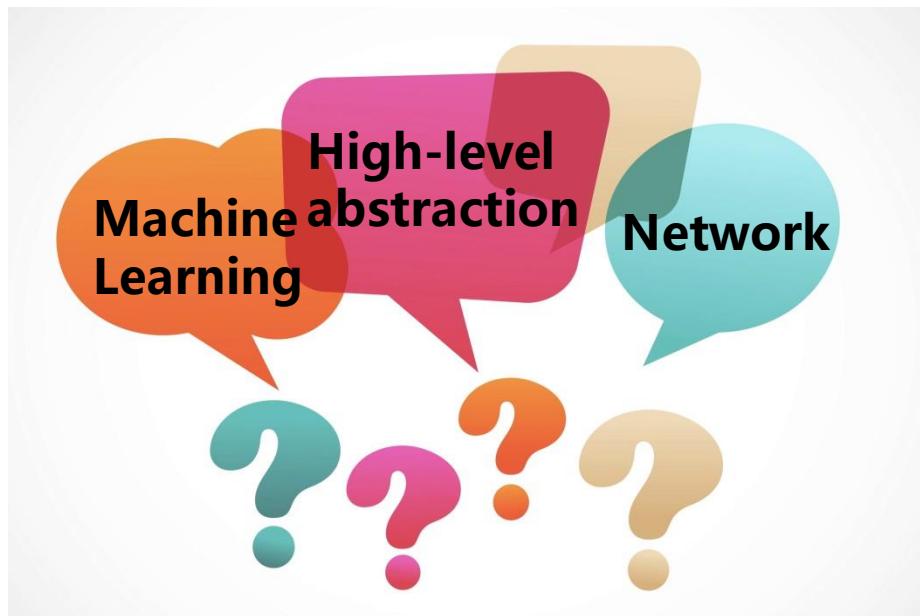
Introduction to deep learning

Sungjoon Choi
[\(sungjoon.choi@cpslab.snu.ac.kr\)](mailto:sungjoon.choi@cpslab.snu.ac.kr)

Deep learning

Wikipedia says:

“Deep learning is a branch of machine learning based on a set of algorithms that attempt to model high-level abstractions in data by using multiple processing layers, with complex structures or otherwise, composed of multiple non-linear transformations.”



Deep learning

Artificial Intelligence

Machine Learning

Deep learning

MLP

CNN

RNN

Is it brand new?



Neural Nets

McCulloch & Pitt 1943

Perceptron

Rosenblatt 1958

RNN

Grossberg 1973

CNN

Fukushima 1979

RBM

Hinton 1999

DBN

Hinton 2006

D-AE

Vincent 2008

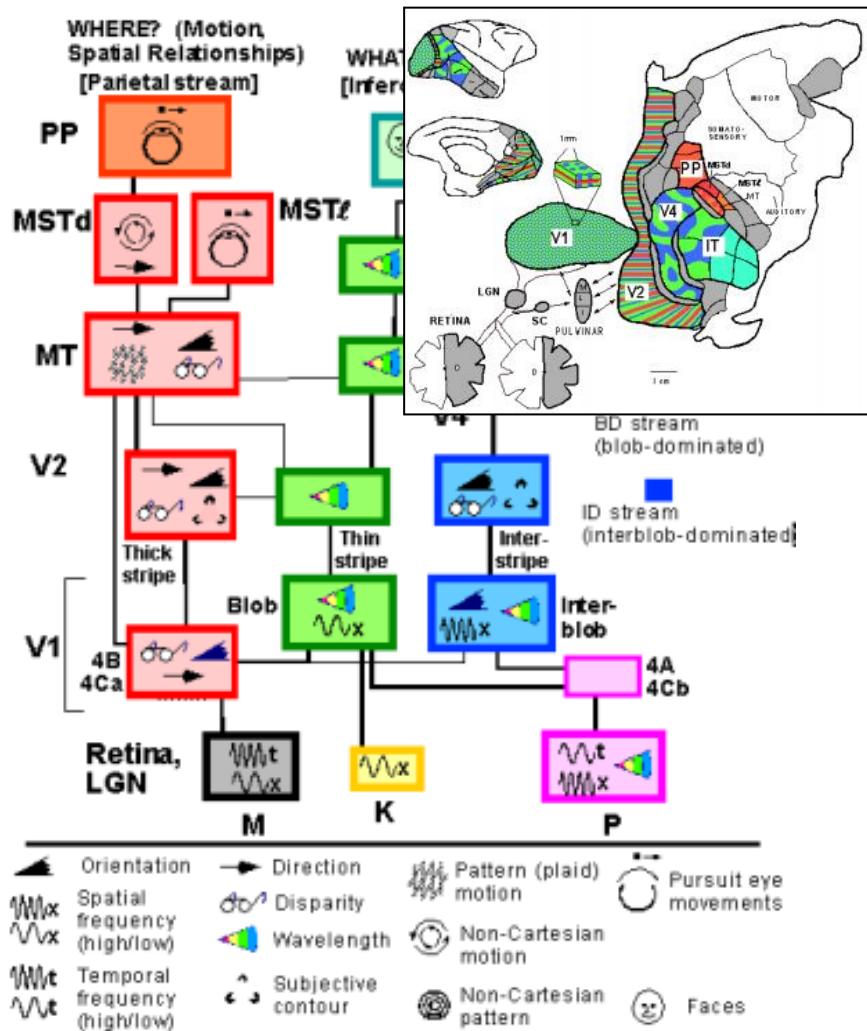
AlexNet

Alex 2012

GoogLeNet

Szegedy 2015

Inspired by nature?



Sure there is some **resemblance!**



L'Avion III de Clement Ader, 1897



But, do not go **TOO FAR.**

Theory

Neural Networks, Vol. 2, pp. 359–366, 1989
Printed in the USA. All rights reserved.

0893-6080/89 \$3.00 + .00
Copyright © 1989 Pergamon Press plc

English

There is a single hidden layer feedforward network that approximates any measurable function to any desired degree of accuracy on some compact set K .

Math

For every function g in M^r there is a compact subset K of R^r and an $f \in \sum^r(\Psi)$ such that for any $\epsilon > 0$ we have $\mu(K) < 1 - \epsilon$ and for every $X \in K$ we have

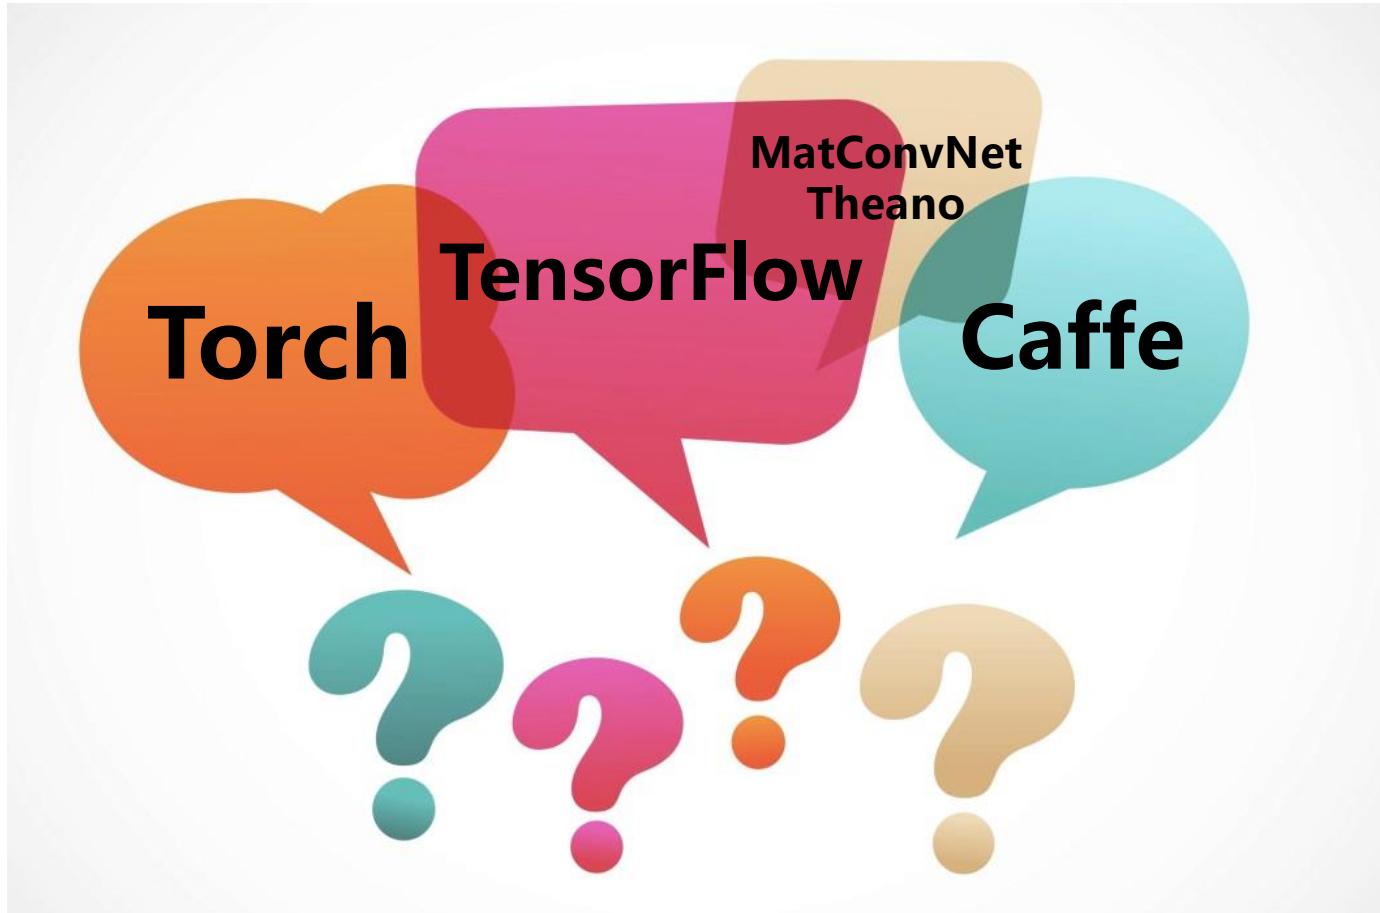
$|f(x) - g(x)| < \epsilon$, regardless of Ψ , r , or μ .

hidden units are available. In this sense, multilayer feedforward networks are a class of universal approximators.

Keywords—Feedforward networks, Universal approximation, Mapping networks, Network representation capability, Stone-Weierstrass Theorem, Squashing functions, Sigma-Pi networks, Back-propagation networks.

Cited > 12,000

Deep learning tools?



White to black

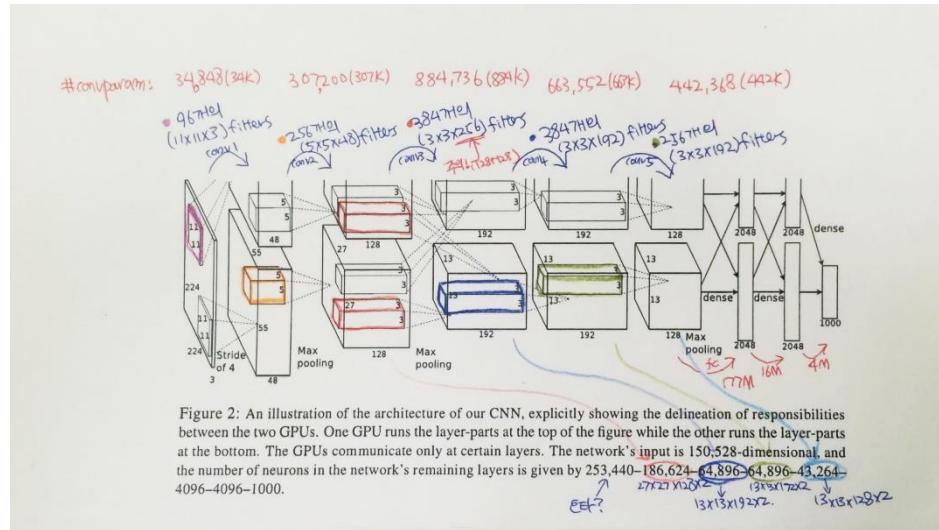


Course Contents

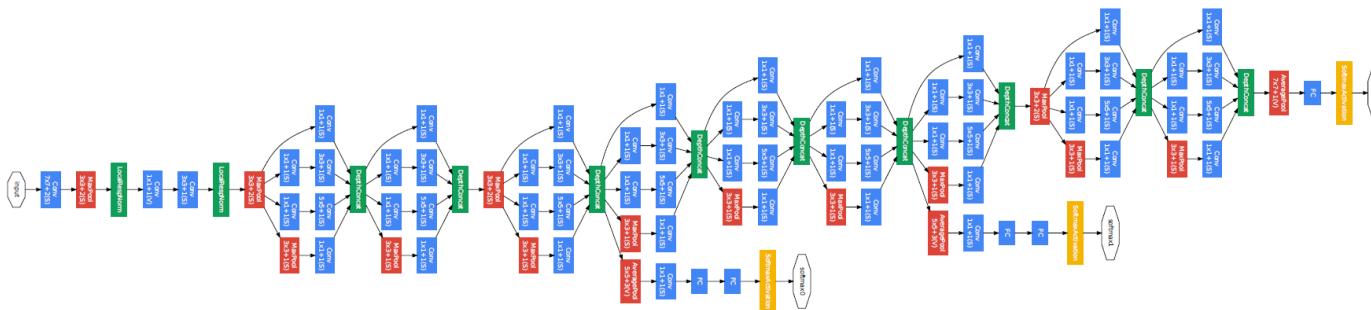
Sungjoon Choi
(sungjoon.choi@cplab.snu.ac.kr)

Convolutional neural network

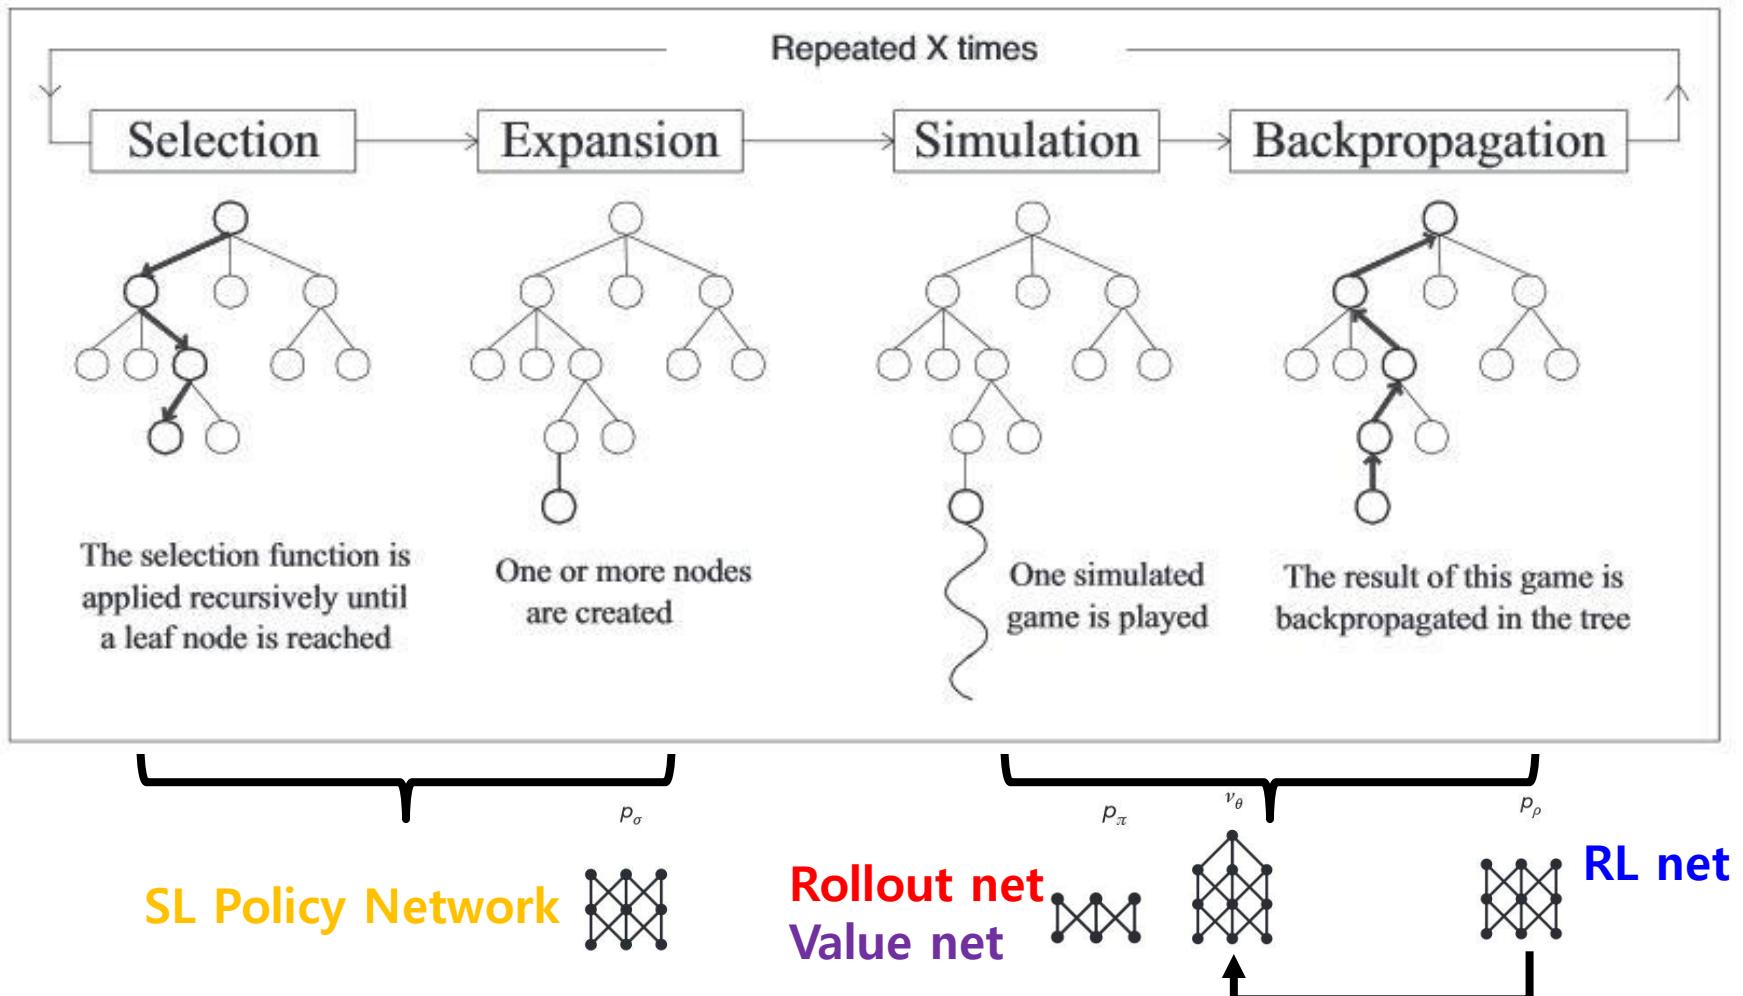
AlexNet



GoogleLeNet



Algorithms behind AlphaGo

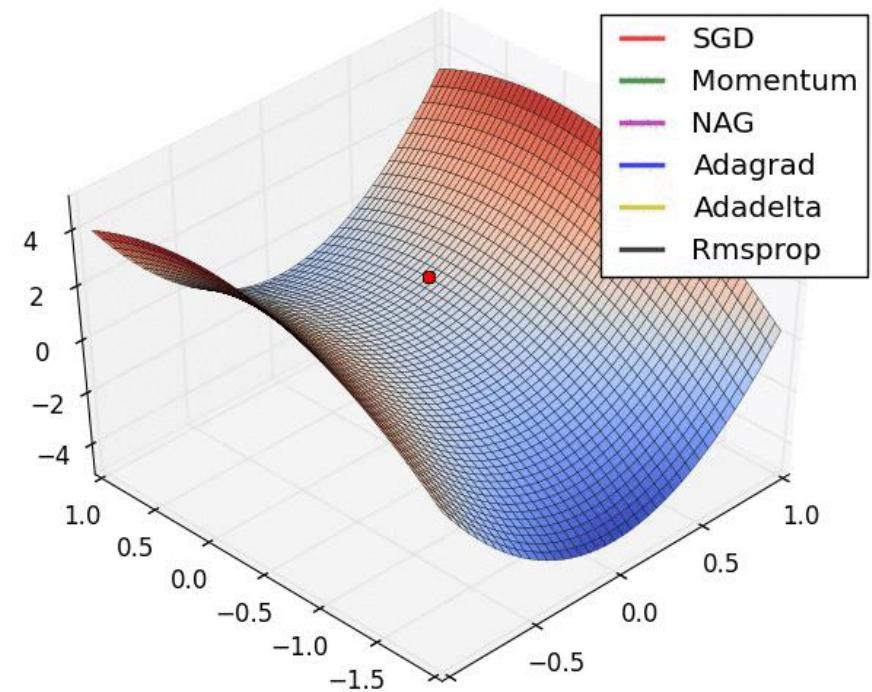
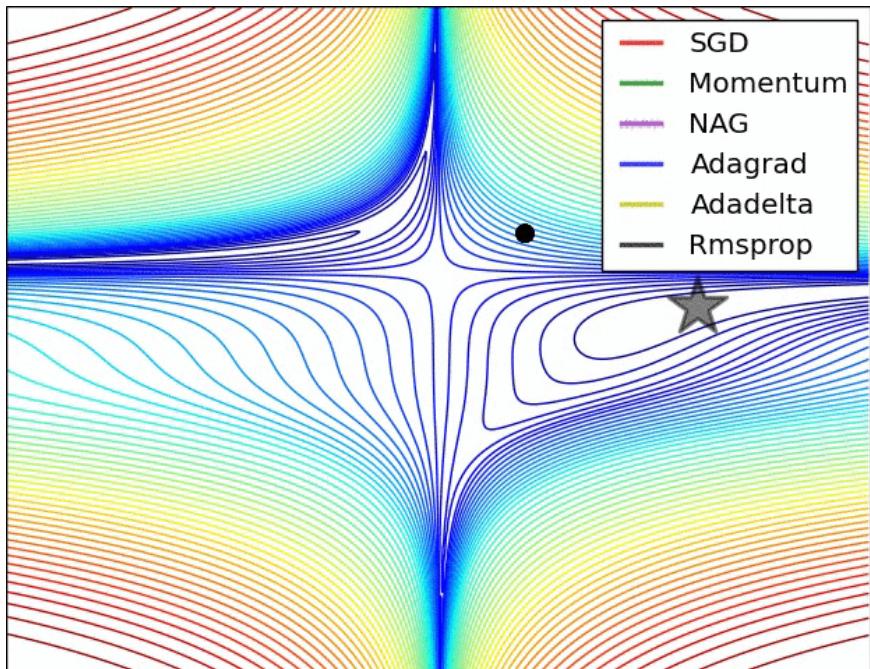


Regularization



Main purpose is to avoid “**OverFitting**”

Optimization Methods



Restricted Boltzmann machine

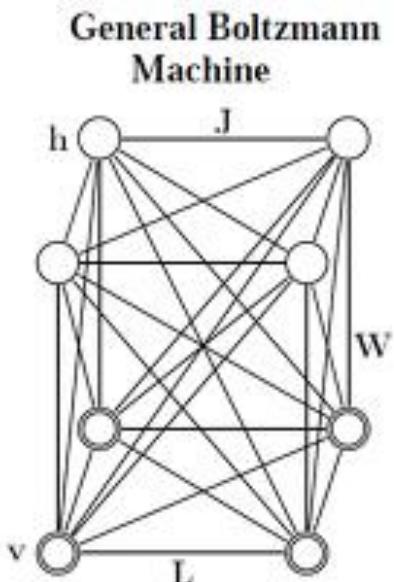


Figure 1: **Left:** A general Boltzmann machine represents a vector of stochastic binary variables. The top layer represents a vector of "latent" variables. The bottom layer represents a vector of "observed" or "visible" variables. **Right:** A restricted Boltzmann machine has no visible-to-visible connections and no hidden-to-hidden connections.

Restricted Boltzmann Machine

Restricted Boltzmann Machine

We don't know h

$$\frac{\partial}{\partial \theta} (-\log P(v)) = E\left[\frac{\partial}{\partial \theta} E(v, h) | v\right] - E\left[\frac{\partial}{\partial \theta} E(v, h)\right]$$

* $E(v, h | \theta) = v^T \underbrace{W h}_{\theta} + \underbrace{b^T v}_{\theta} + \underbrace{a^T h}_{\theta}$ we know
we don't know both of v and h .

positive phase negative phase

Gibbs sampled h

Gibbs sampled v

Known visible nodes

Contrastive Divergence (CD) is used.

$$\frac{\partial}{\partial w_{ij}} E(v, h | \theta) = v_i h_j$$

$$\frac{\partial}{\partial b_i} E(v, h | \theta) = v_i$$

$$\frac{\partial}{\partial a_j} E(v, h | \theta) = h_j$$

T W Done

Semantic Segmentation

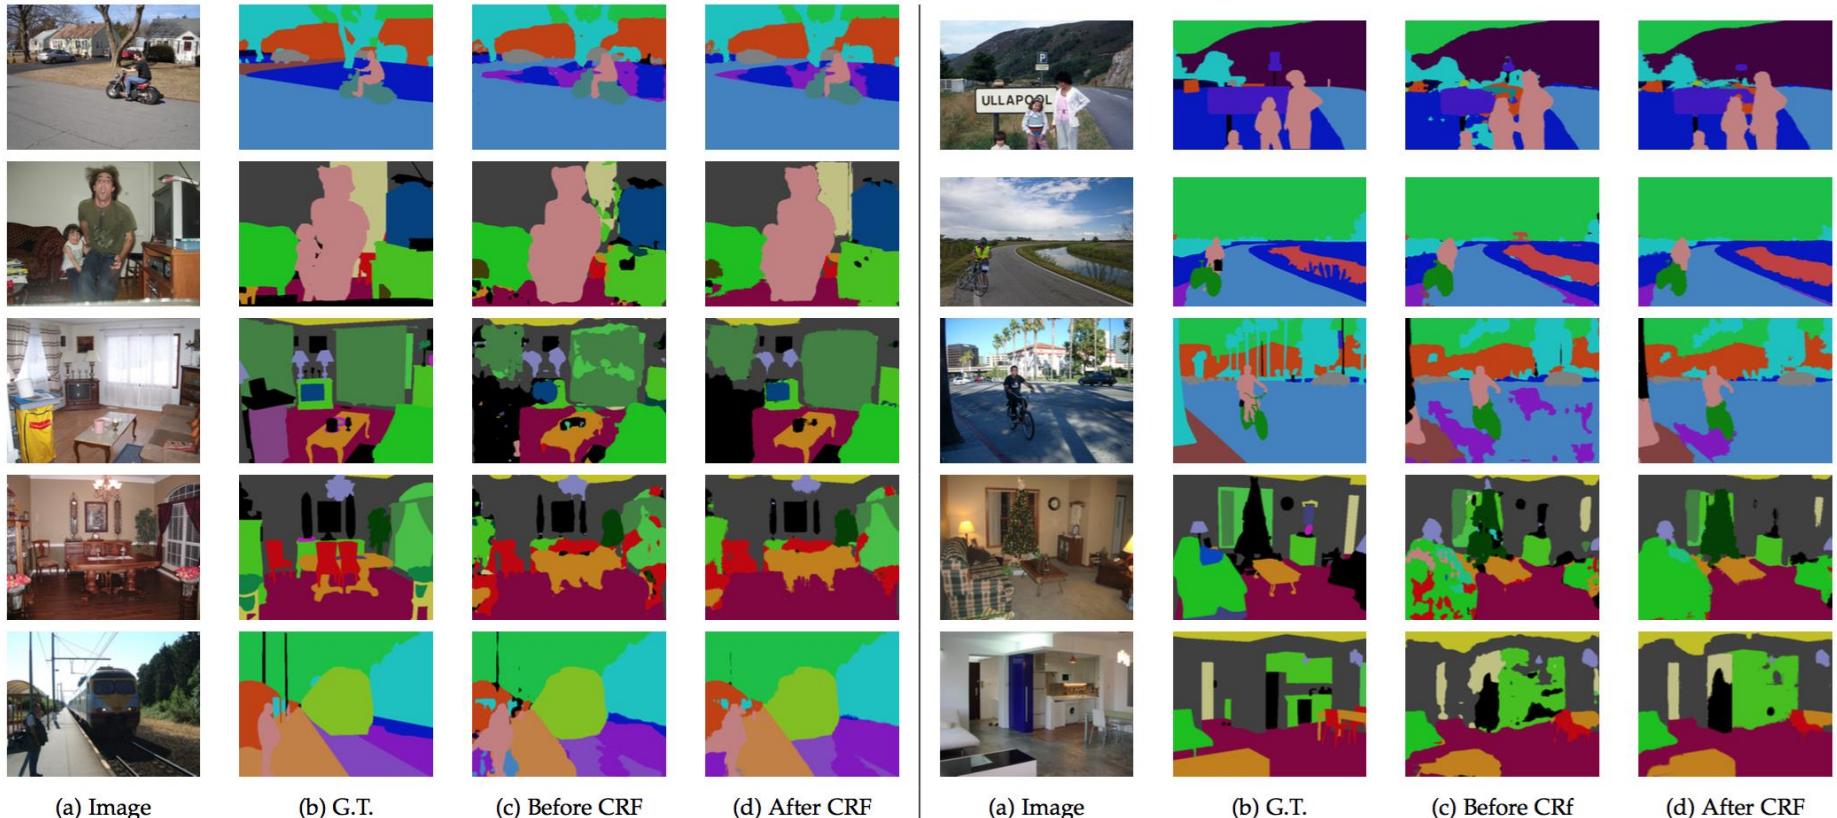
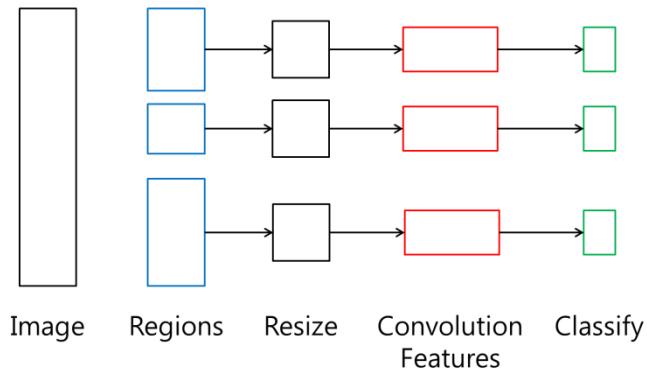
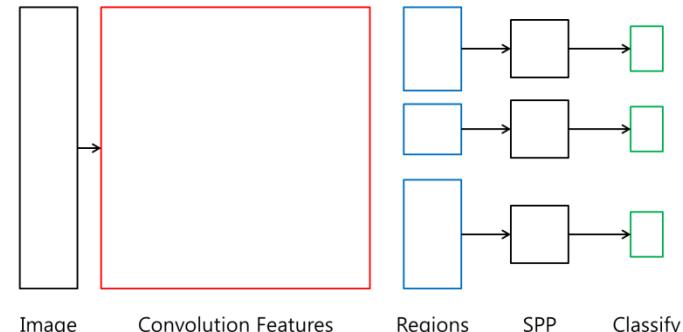


Fig. 11: PASCAL-Context results. Input image, ground-truth, and our DeepLab results before/after CRF.

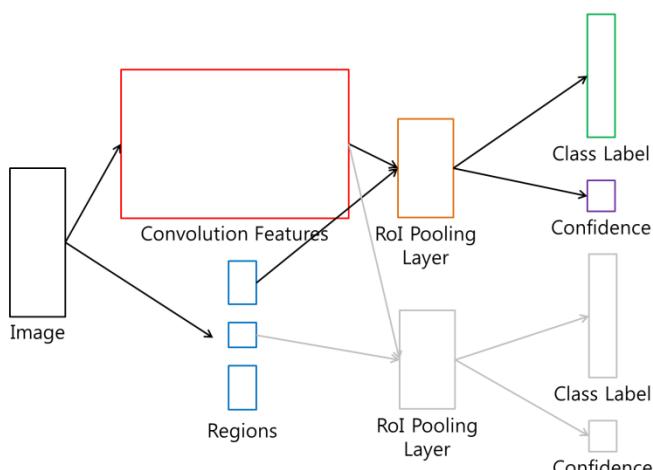
Detection Methods



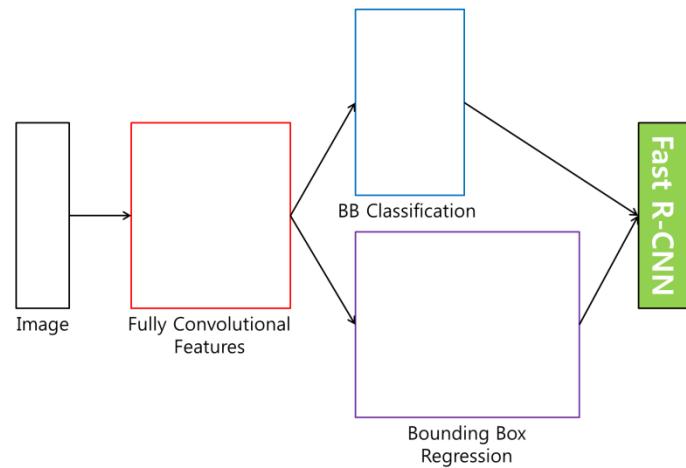
R-CNN



SPP net

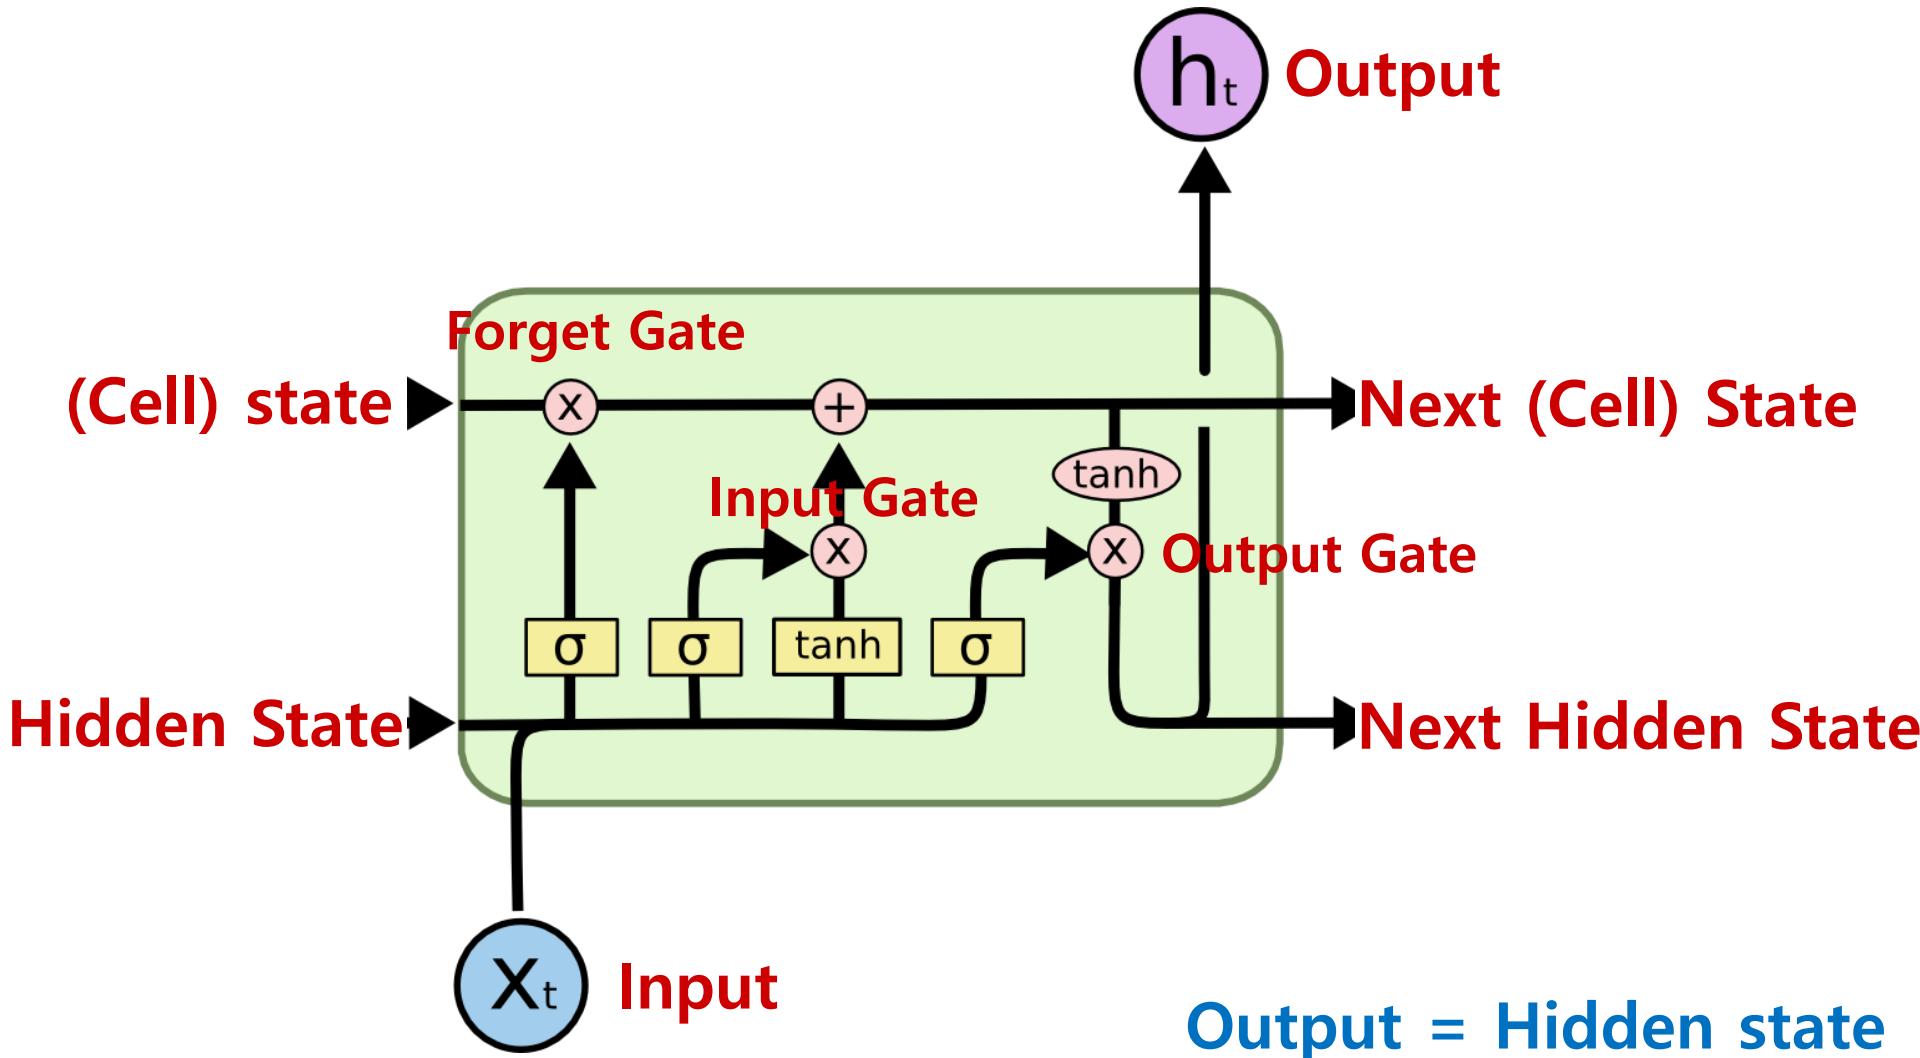


Fast R-CNN

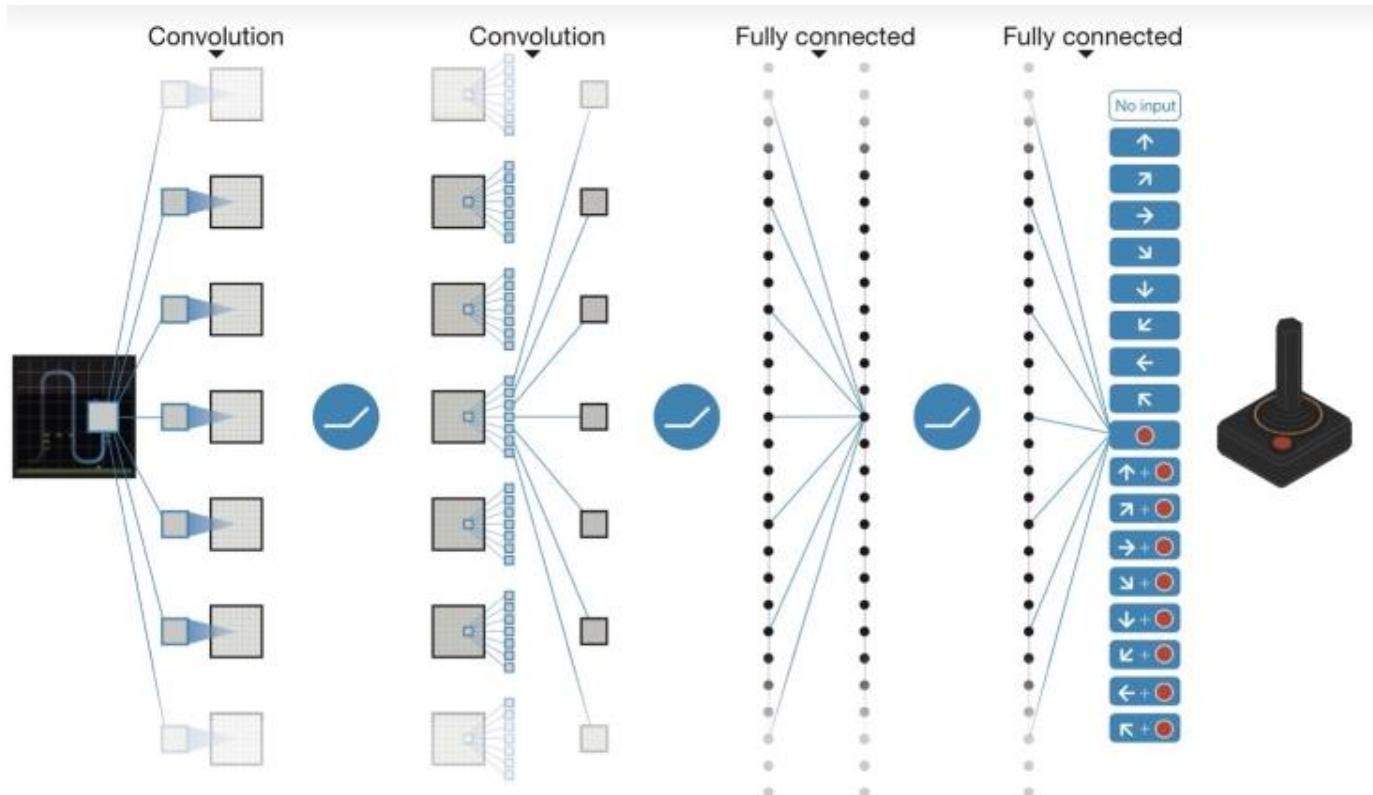


Faster R-CNN

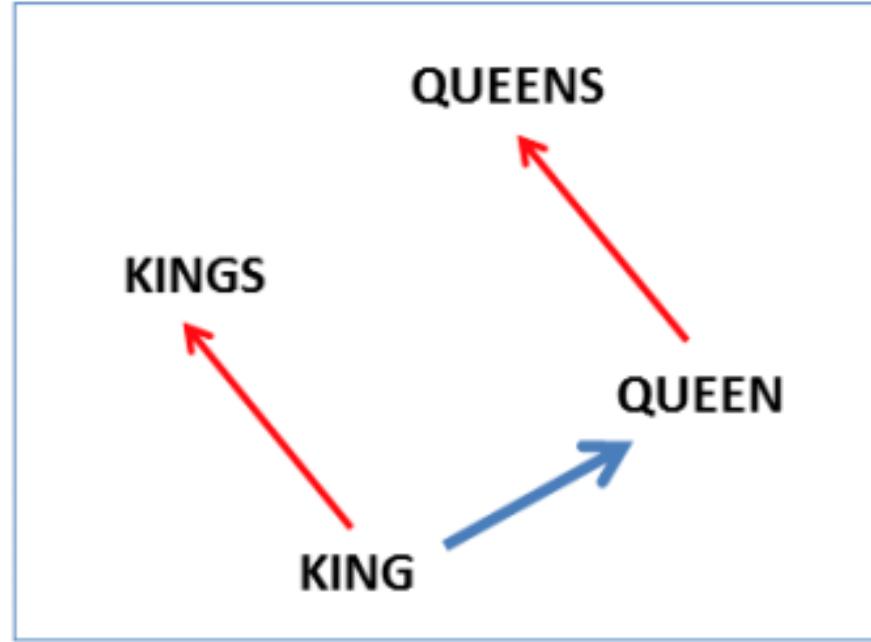
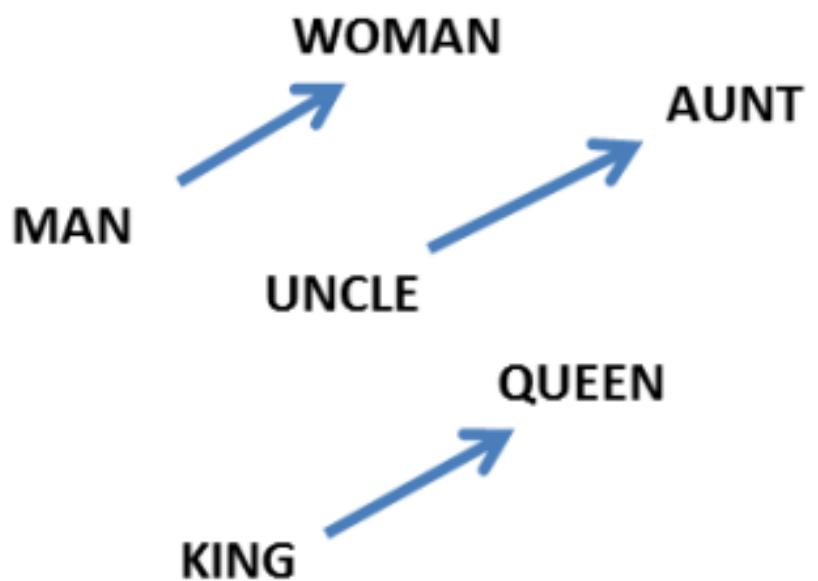
RNN - LSTM



Deep Reinforcement Learning

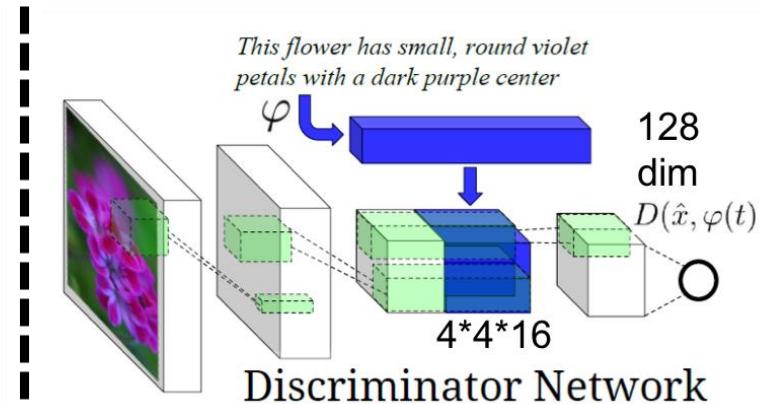
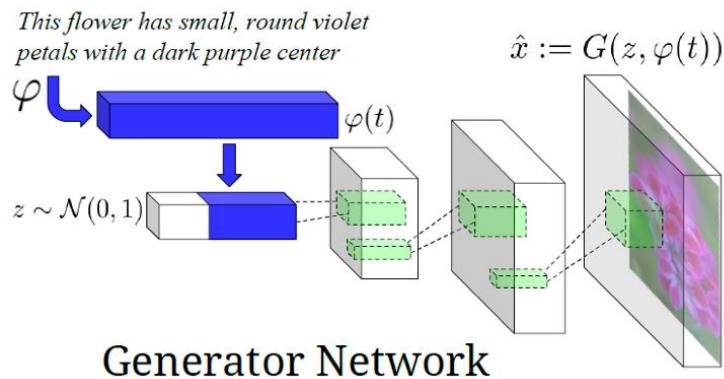
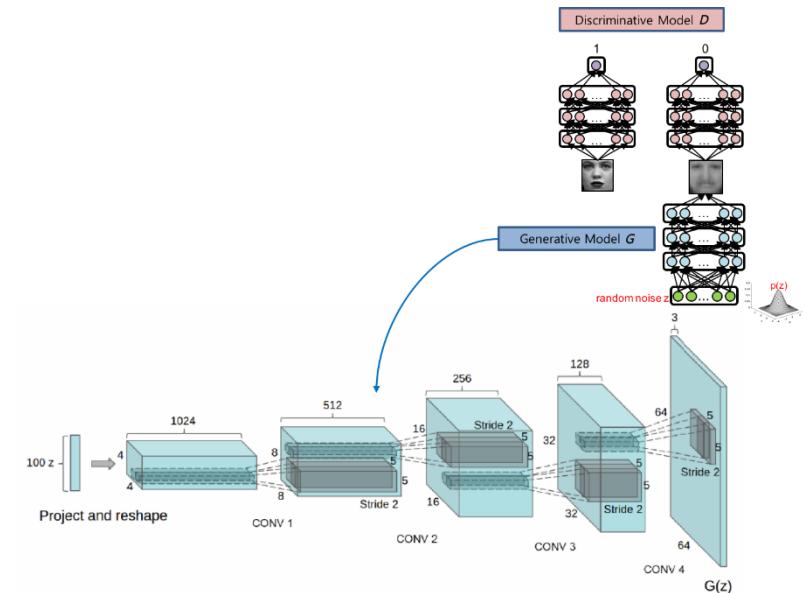
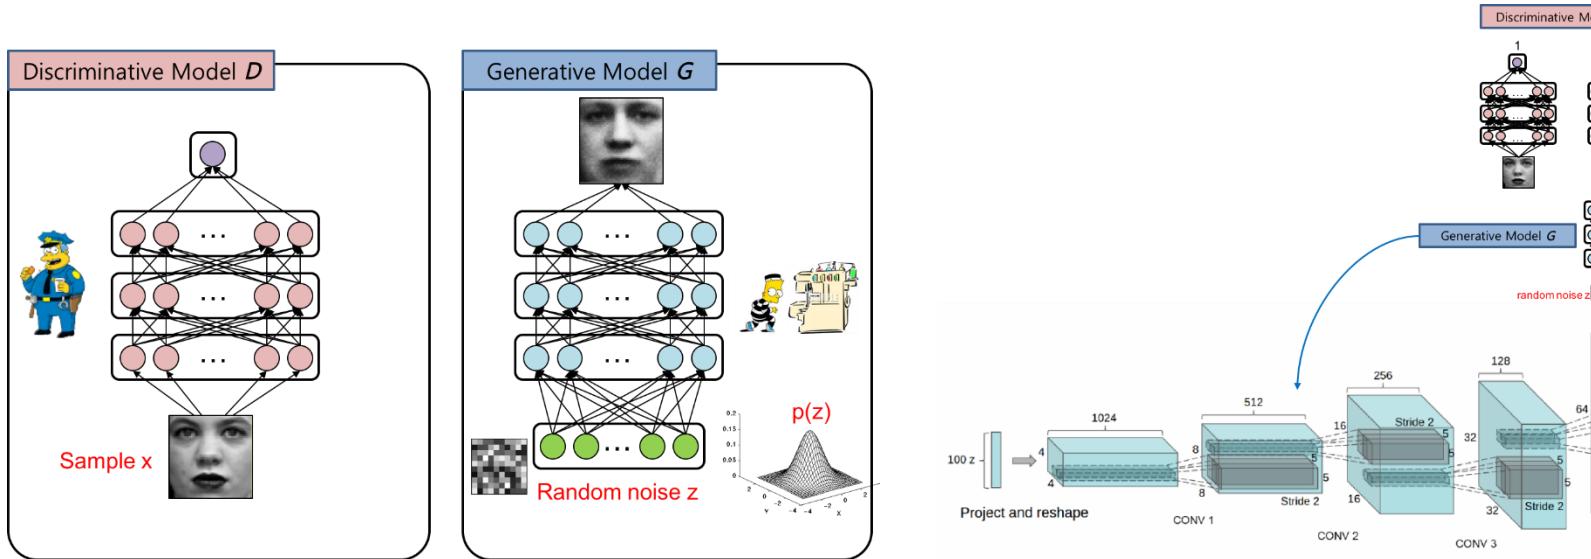


Word Embedding



(Mikolov et al., NAACL HLT, 2013)

Generative Adversarial Networks



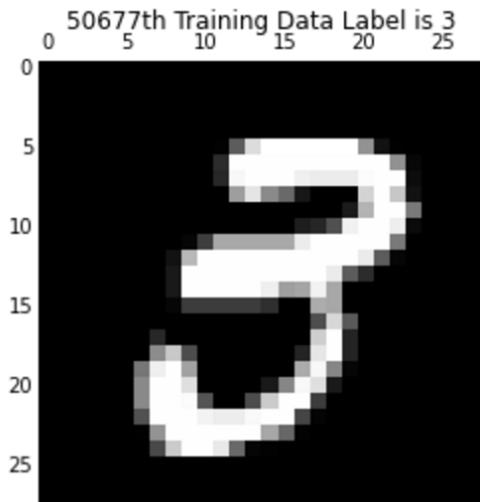
+ Alpha

Using MNIST

```
# How does the training data look like?
print ("How does the data look like?")
nsample = 1
randidx = np.random.randint(trainimg.shape[0], size=nsample)

for i in randidx:
    curr_img = np.reshape(trainimg[i, :], (28, 28)) # 28 by 28 matrix
    curr_label = np.argmax(trainlabel[i, :]) # Label
    plt.matshow(curr_img, cmap=plt.get_cmap('gray'))
    plt.title("'" + str(i) + "th Training Data "
              + "Label is " + str(curr_label))
    print ("'" + str(i) + "th Training Data "
          + "Label is " + str(curr_label))
```

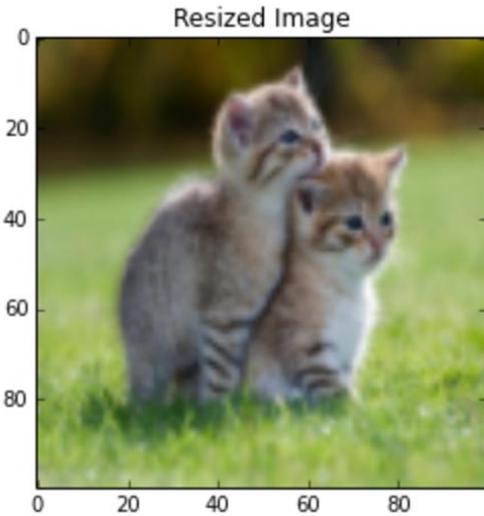
How does the data look like?
50677th Training Data Label is 3



Basic Image Handling

```
# Resize
catsmall = imresize(cat, [100, 100, 3])
print_type(shape(catsmall))
# Plot
plt.figure(1)
plt.imshow(catsmall)
plt.title("Resized Image")
plt.draw()
```

Type is <type 'numpy.ndarray'>
Shape is (100, 100, 3)

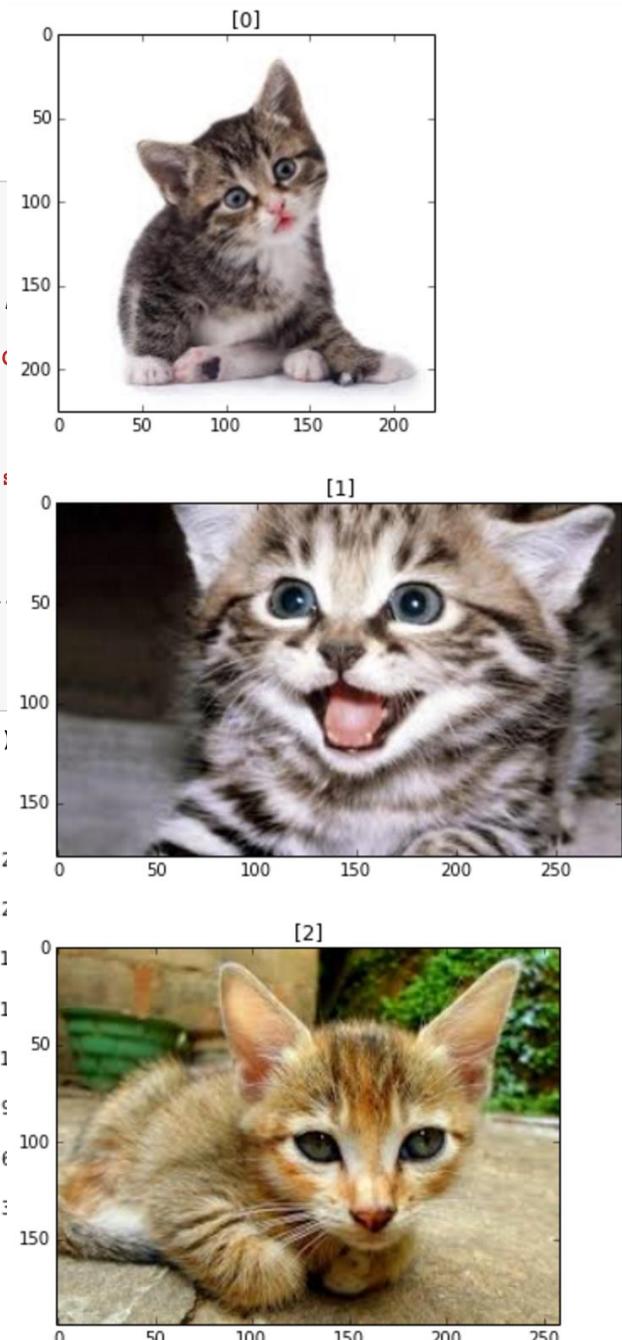
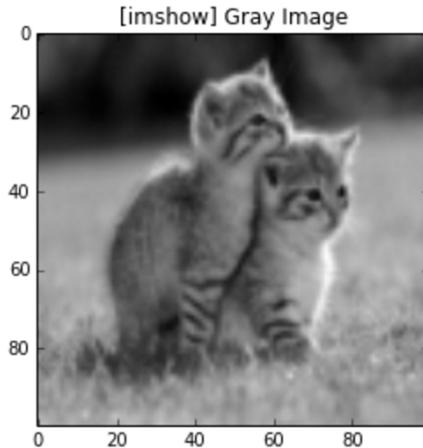


```
# Grayscale
def rgb2gray(rgb):
    if len(rgb.shape) is 3:
        return np.dot(rgb[...,:3], [0.299, 0.587, 0.144])
    else:
        print ("Current Image is not in RGB format")
        return rgb
catsmallgray = rgb2gray(catsmall)

print ("size of catsmallgray is %s" % catsmallgray.shape)
print ("type of catsmallgray is", type(catsmallgray))

plt.figure(2)
plt.imshow(catsmallgray, cmap=plt.cm.Greys)
plt.title("[imshow] Gray Image")
plt.colorbar()
plt.draw()
```

size of catsmallgray is (100, 100)
(type of catsmallgray is, <type



Denoising Auto-Encoder

```
# tf Graph input
x = tf.placeholder("float", [None, n_input])
y = tf.placeholder("float", [None, n_output])
dropout_keep_prob = tf.placeholder("float")
# Store layers weight & bias
weights = {
    'h1': tf.Variable(tf.random_normal([n_input, n_hidden_1])),
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_hidden_2, n_output]))
}
biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_output]))
}
def encoder(x, weights, biases):
    layer_1 = tf.nn.sigmoid(tf.matmul(x, weights['h1']) + biases['b1'])
    layer_1out = tf.nn.dropout(layer_1, dropout_keep_prob)
    layer_2 = tf.nn.sigmoid(tf.add(tf.matmul(layer_1out, weights['h2']), biases['b2']))
    layer_2out = tf.nn.dropout(layer_2, dropout_keep_prob)
    return tf.nn.sigmoid(tf.matmul(layer_2out, weights['out']) + biases['out'])

# Build a Denoising Autoencoder
original_image = np.reshape(X, [-1, 28, 28, 1])
input_image = np.reshape(X, [-1, 28, 28, 1])
input_image = np.clip(input_image + np.random.normal(0, 0.1, input_image.shape), 0, 1)
reconstructed_image = encoder(input_image, weights, biases)

# Train the model
optimizer = tf.train.AdamOptimizer(0.001).minimize(cost)
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)
```



The figure displays three 28x28 pixel grayscale images. The first image, labeled 'Original Image', shows a clear handwritten digit '2'. The second image, labeled 'Input Image', is identical to the first but contains scattered white noise pixels. The third image, labeled 'Reconstructed Image', is the output of the autoencoder and appears nearly identical to the 'Original Image', indicating that the noise has been effectively removed.

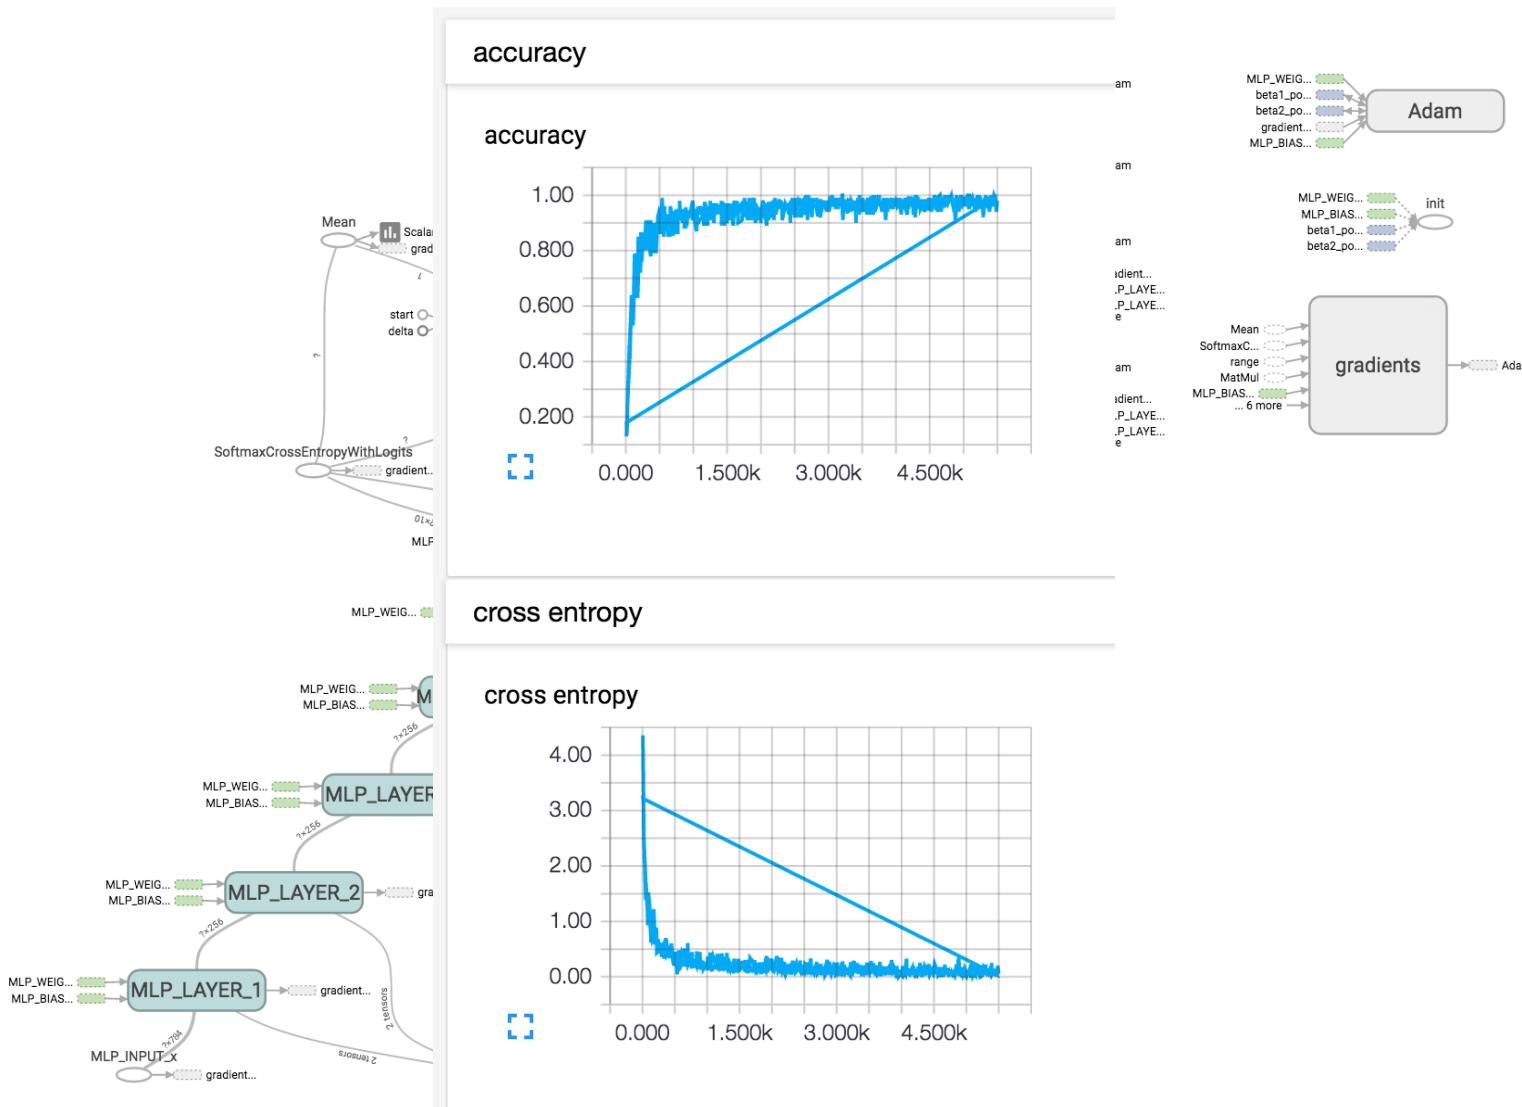
Convolutional Neural Network

```
weights = {
    'wc1': tf.Variable(tf.random_normal([3, 3, 1, 64], stddev=0.1)),
    'wc2': tf.Variable(tf.random_normal([3, 3, 64, 128], stddev=0.1)),
    'wd1': tf.Variable(tf.random_normal([7*7*128, 1024], stddev=0.1)),
    'wd2': tf.Variable(tf.random_normal([1024, n_output], stddev=0.1))
}
biases = {
    'bc1': tf.Variable(tf.random_normal([64], stddev=0.1)),
    'bc2': tf.Variable(tf.random_normal([128], stddev=0.1)),
    'bd1': tf.Variable(tf.random_normal([1024], stddev=0.1)),
    'bd2': tf.Variable(tf.random_normal([n_output], stddev=0.1))
}
def conv_basic(_input, _w, _b, _keepratio):
    # Input
    _input_r = tf.reshape(_input, shape=[-1, 28, 28, 1])
    # Conv1
# Compute test accuracy
test_acc = sess.run(accr, feed_dict={x: testimg, y: testlabel, keepratio:1.})
print (" Test accuracy: %.3f" % (test_acc))
```

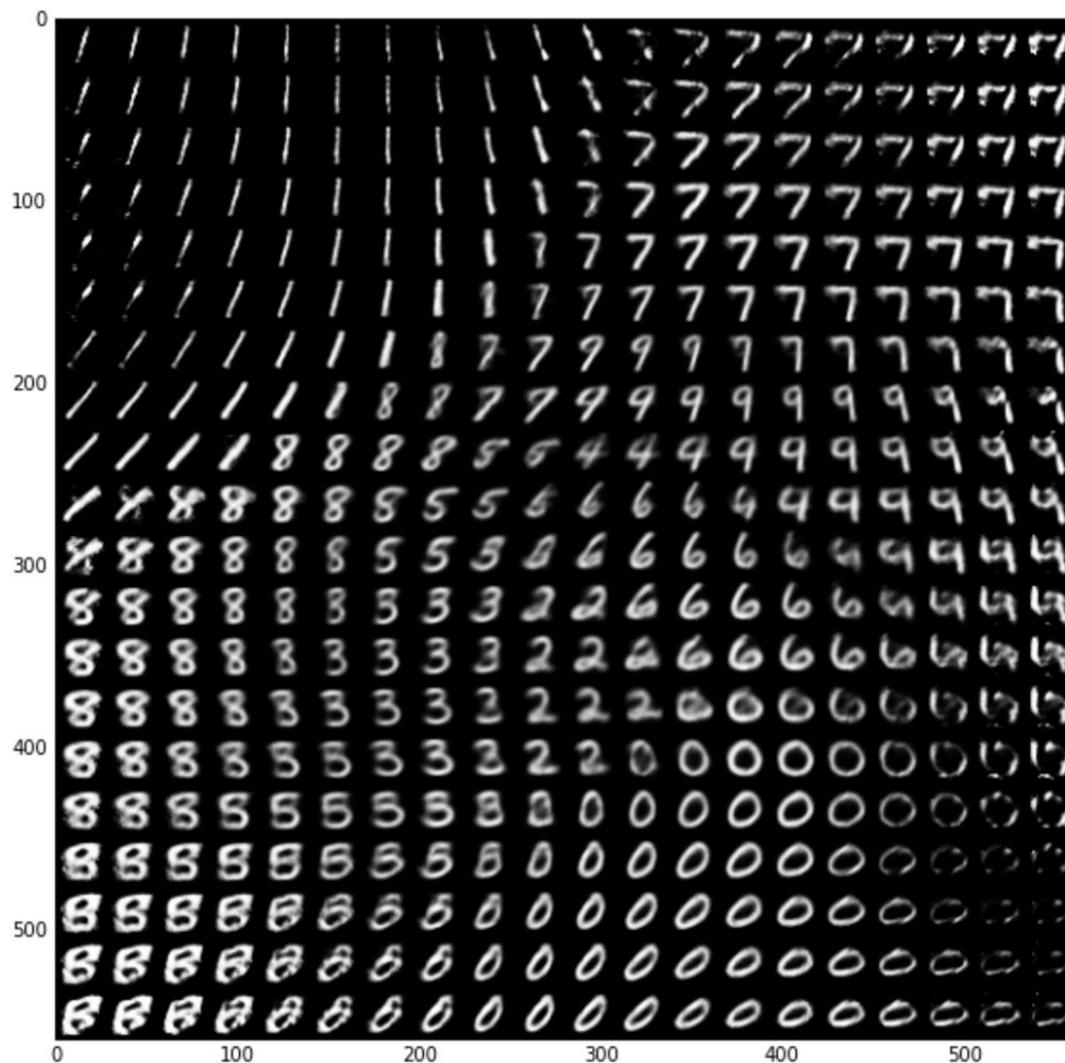
Test accuracy: 0.994

```
# Vectorize
_dense1 = tf.reshape(_pool_dr2, [-1, _w['wd1'].get_shape().as_list()[0]])
# Fc1
_fc1 = tf.nn.relu(tf.add(tf.matmul(_dense1, _w['wd1']), _b['bd1']))
_fc_dr1 = tf.nn.dropout(_fc1, _keepratio)
# Fc2
_out = tf.add(tf.matmul(_fc_dr1, _w['wd2']), _b['bd2'])
# Return everything
out = {
    'input_r': _input_r,
    'conv1': _conv1,
    'pool1': _pool1,
    'pool1_dr1': _pool_dr1,
    'conv2': _conv2,
    'pool2': _pool2,
    'pool_dr2': _pool_dr2,
    'dense1': _dense1,
    'fc1': _fc1,
    'fc_dr1': _fc_dr1,
    'out': _out
}
return out
```

Tensor Board



Variational Autoencoder (VAE)

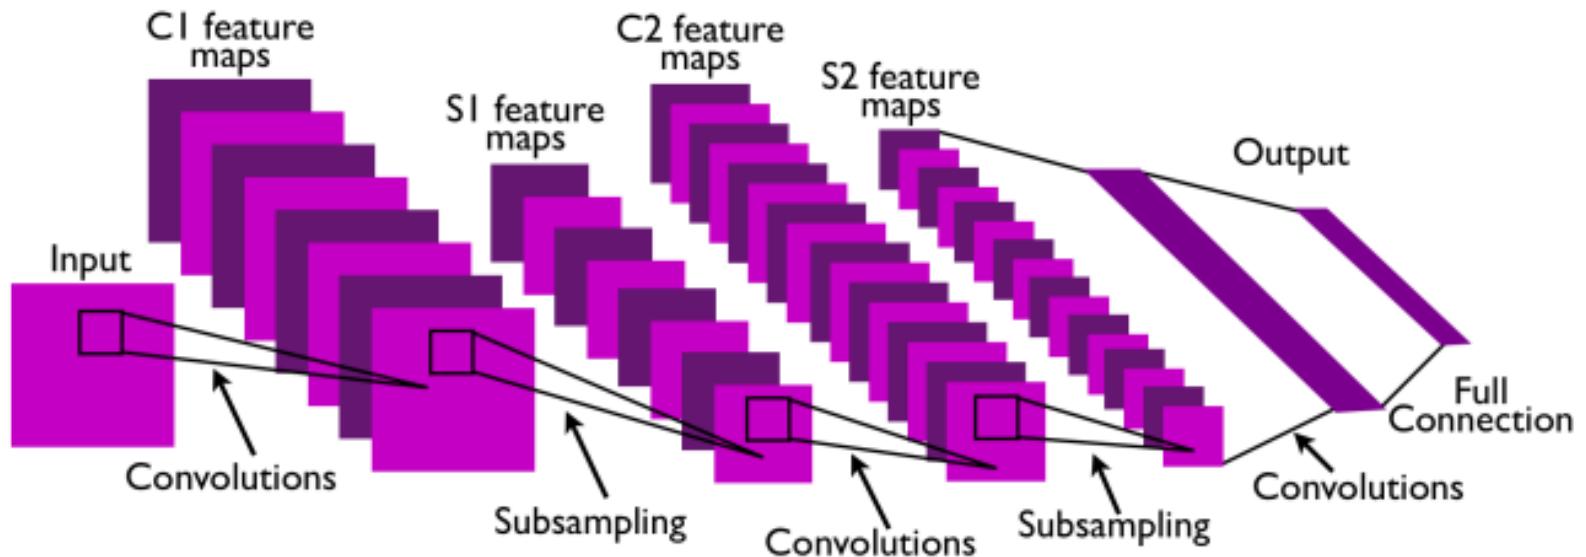


Introduction to CNNs

Sungjoon Choi
[\(sungjoon.choi@cpslab.snu.ac.kr\)](mailto:sungjoon.choi@cpslab.snu.ac.kr)

CNN

Convolutional Neural Network



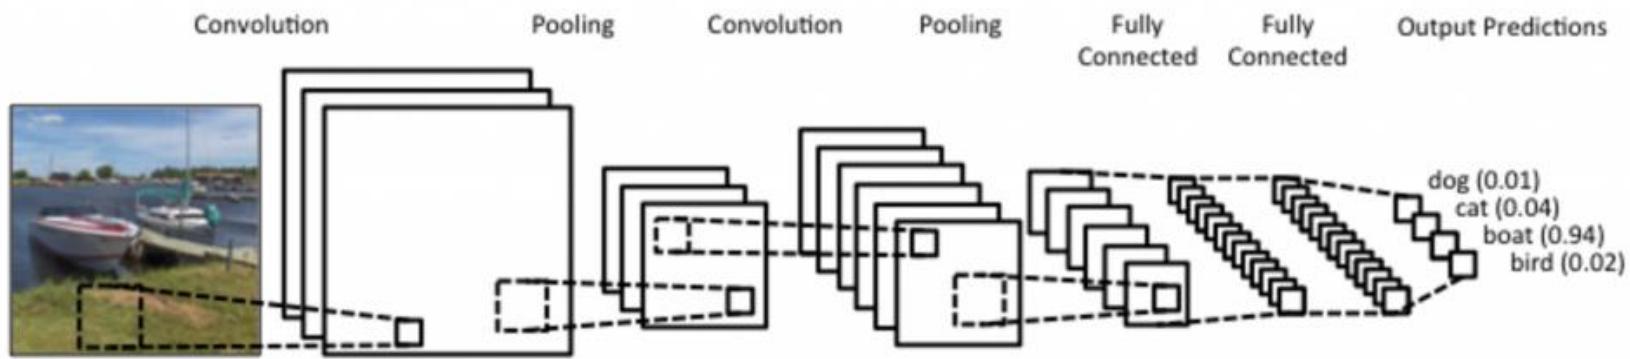
This is pretty much **everything** about the convolutional neural network.

Convolution + Subsampling + Full Connection

CNN

CNNs are basically layers of **convolutions** followed by **subsampling** and **fully connected layers**.

Intuitively speaking, **convolutions** and **subsampling** layers works as feature extraction layers while a **fully connected layer** classifies which category current input belongs to using extracted features.



Why so powerful?



Local Invariance

Loosely speaking, as the convolution filters are '**sliding**' over the input image, the exact location of the object we want to find does not matter much.

Compositionality

There is a hierarchy in CNNs. It is GOOD!

Convolution

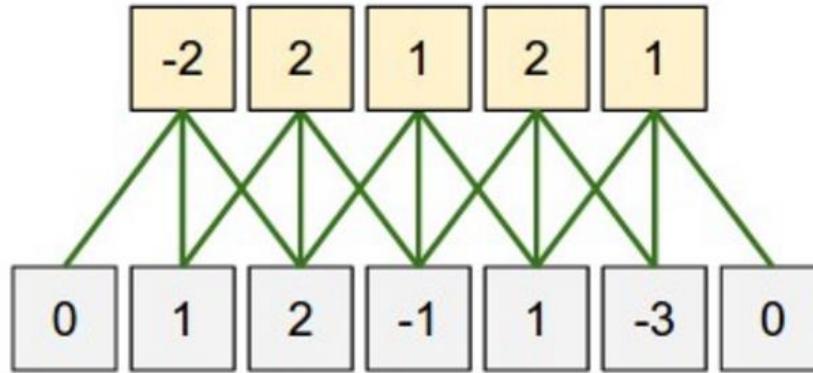
1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Zero-padding



What is the size of the input? $n_{in} = 5$

What is the size of the output? $n_{out} = 5$

What is the size of the filter? $n_{filter} = 3$

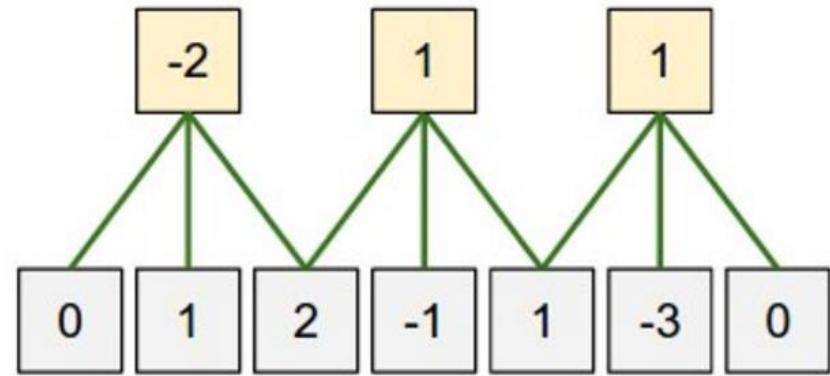
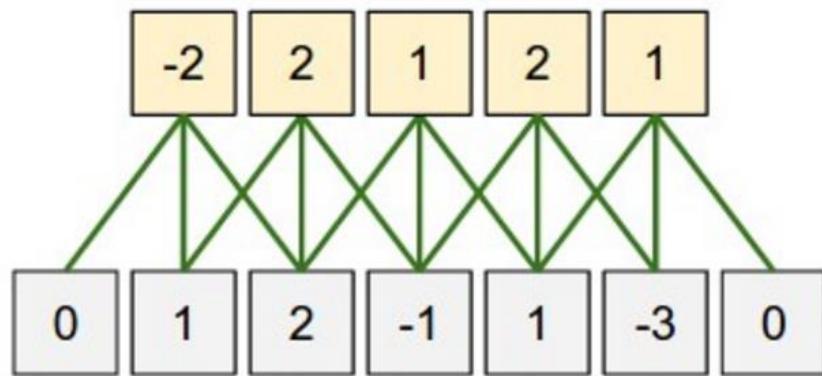
What is the size of the zero-padding? $n_{padding} = 1$

$$n_{out} = (n_{in} + 2 * n_{padding} - n_{filter}) + 1 \quad 5 = (5 + 2 * 1 - 3) + 1$$

Stride



Stride



(Left) Stride size: 1

(Right) Stride size: 2

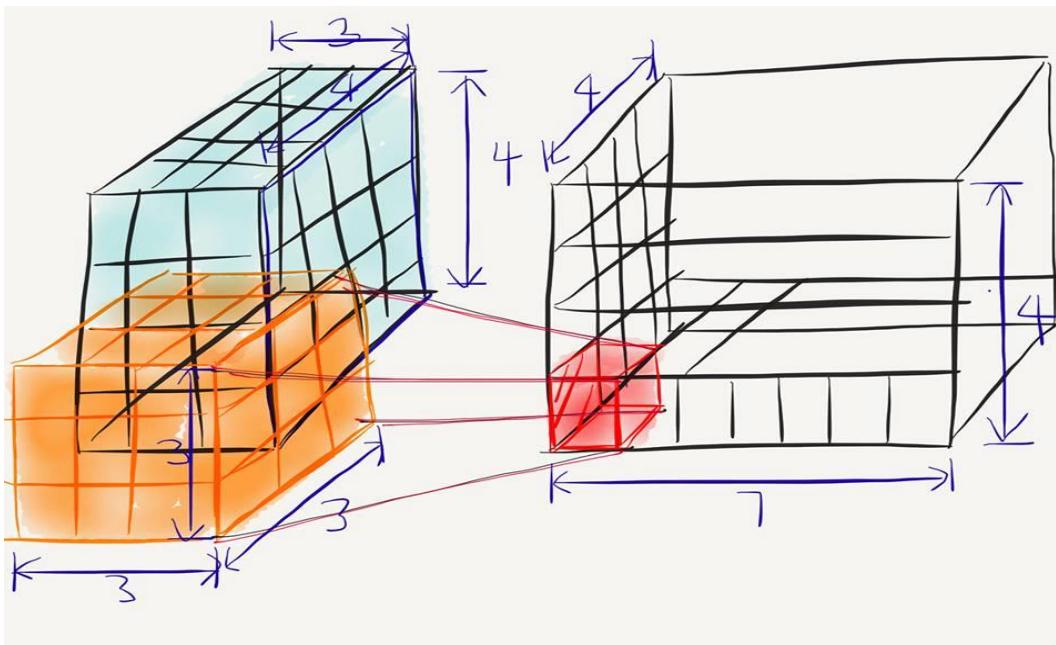
If stride size equals the filter size, there will be **no overlapping**.

Conv2D

```
tf.nn.conv2d(input, filter, strides, padding,  
use_cudnn_on_gpu=None, name=None)
```

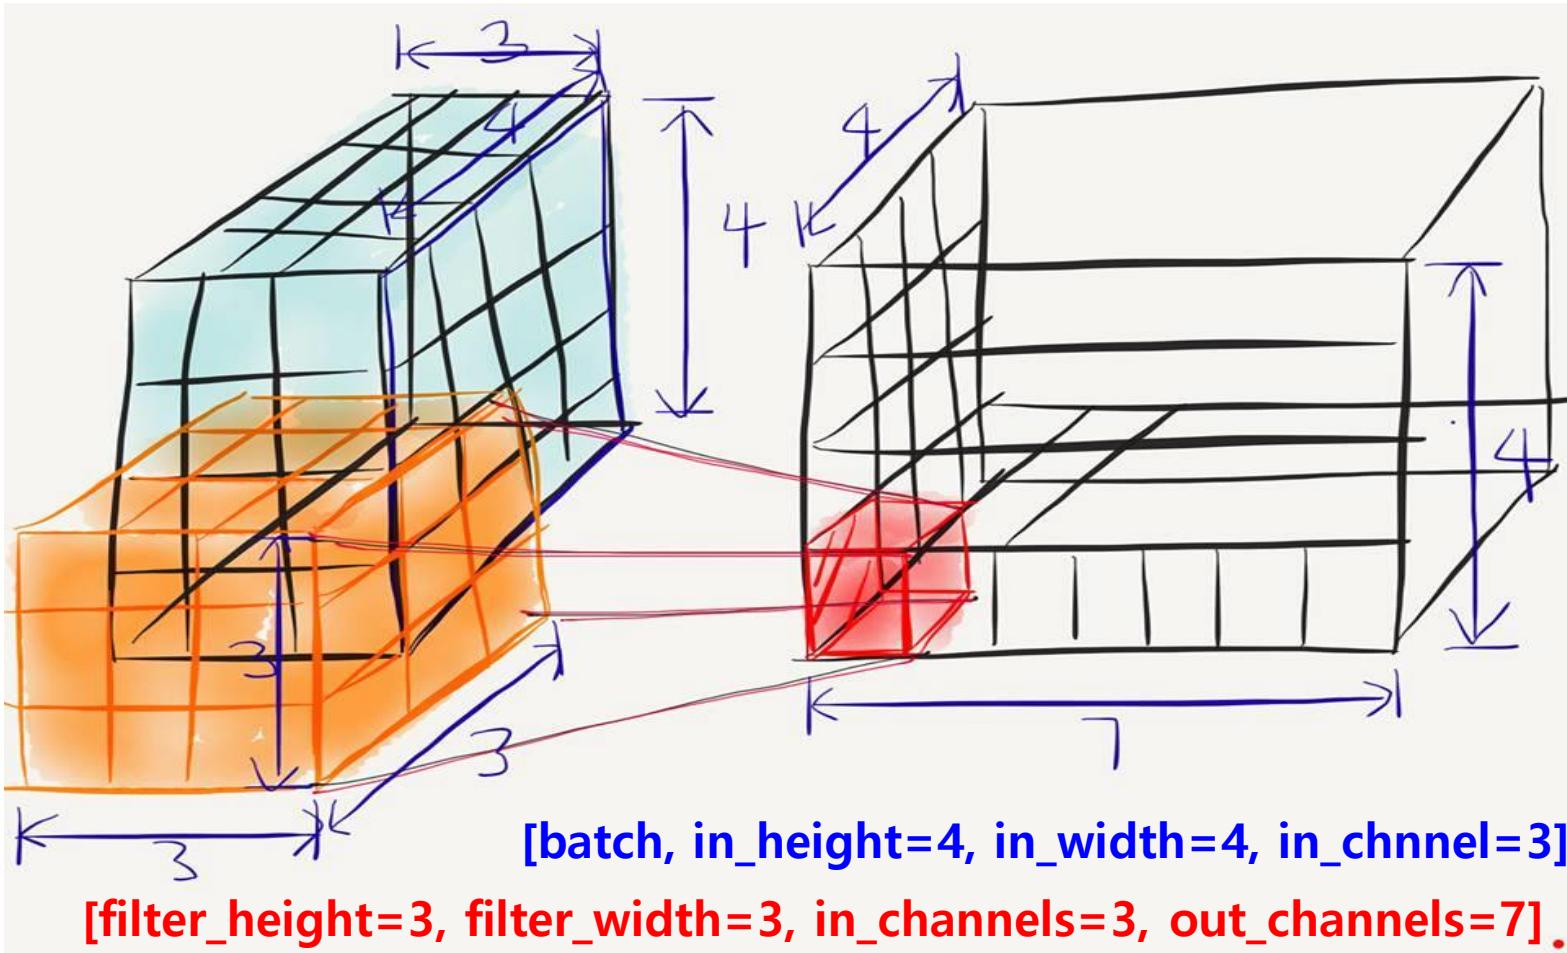
Computes a 2-D convolution given 4-D input and filter tensors.

Given an input tensor of shape **[batch, in_height, in_width, in_channel]** and a filter / kernel tensor of shape **[filter_height, filter_width, in_channels, out_channels]** this op performs the following:



**[batch, in_height=4,
in_width=4, in_channel=3]**
**[filter_height=3, filter_width=3,
in_channels=3, out_channels=7]**

Conv2D



What is the **number of parameters** in this convolution layer?

$$\rightarrow 189 = 3 * 3 * 3 * 7$$

